

Below is a **copy-pasteable, extremely detailed project outline** you can hand to a fresh LLM chat to build the full system. It includes **repo structure**, **Docker Compose**, **file-by-file responsibilities**, **APIs**, **schemas**, **storage**, **Redis logic**, **Kafka topics**, and a **stepwise build plan**.

---

# Project: Internal Developer Event Platform (Ingestion + Processing + Admin Ops)

## Goal

Build an internal, developer-facing event platform with:

- **Fast ingestion API** (Go) with **API-key auth, rate limiting, idempotency**
  - **Async pipeline via Kafka** (`events.raw`, `events.dlq`)
  - **Processor** (Java Spring Boot) that validates/enriches events and stores them in **Postgres**
  - **DLQ routing for poison messages**
  - **Admin API** (can live in processor service) exposing operational metrics
  - **React admin dashboard** consuming Admin API
  - **Python load generator** simulating producers (normal, duplicates, malformed, multi-tenant)
  - Everything runs locally via **Docker Compose** for infra (Kafka/Redis/Postgres)
- 

## 1) Repo Structure (Monorepo)

pulse-event-platform/

  README.md

  .gitignore

  infra/

    docker-compose.yml

    kafka-init/

      create-topics.sh

  ingest-go/

```
go.mod
cmd/
  ingest/
    main.go
  internal/
    config/
      config.go
  http/
    router.go
    middleware_auth.go
    middleware_ratelimit.go
    handler_events.go
  models.go
  responses.go
redis/
  client.go
  ratelimit.go
  idempotency.go
kafka/
  producer.go
metrics/
  metrics.go
util/
  time.go
  uuid.go
Dockerfile (optional later)
```

```
processor-java/
  pom.xml
  src/main/java/com/yourorg/processor/
    ProcessorApplication.java
    config/
      KafkaConfig.java
      PostgresConfig.java
      MetricsConfig.java
    kafka/
      RawEventConsumer.java
      DlqProducer.java
      EventPublisher.java (optional)
  service/
    EventProcessorService.java
    ValidationService.java
    EnrichmentService.java
  db/
```

```
EventEntity.java  
EventRepository.java  
Migrations.md (notes)  
api/  
    AdminController.java  
    DtoModels.java  
observability/  
    LagService.java  
    HealthIndicators.java  
util/  
    JsonUtil.java  
src/main/resources/  
    application.yml  
db/migration/  
    V1__init.sql  
Dockerfile (optional later)
```

```
dashboard/  
package.json  
next.config.js  
src/  
    main.tsx  
    App.tsx  
    api/  
        client.ts  
        types.ts  
    pages/  
        Overview.tsx  
        Pipeline.tsx  
        Dlq.tsx  
        Search.tsx  
    components/  
        MetricCard.tsx  
        Chart.tsx (optional)  
        Table.tsx  
    styles/  
        index.css
```

```
loadgen/  
requirements.txt  
loadgen.py  
scenarios/  
    baseline.json  
    spike.json
```

**Key point:** For MVP, you can run services directly on Windows terminals (Go/Java/React/Python) while Docker runs Kafka/Redis/Postgres.

---

## 2) Infrastructure: Docker Compose (Kafka + Redis + Postgres)

Create: `infra/docker-compose.yml`

```
services:
  postgres:
    image: postgres:16
    container_name: ep_postgres
    environment:
      POSTGRES_USER: ep_user
      POSTGRES_PASSWORD: ep_pass
      POSTGRES_DB: event_platform
    ports:
      - "5432:5432"
    volumes:
      - ep_postgres_data:/var/lib/postgresql/data
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U ep_user -d event_platform"]
      interval: 5s
      timeout: 5s
      retries: 10

  redis:
    image: redis:7
    container_name: ep_redis
    ports:
      - "6379:6379"
    command: ["redis-server", "--appendonly", "yes"]
    volumes:
      - ep_redis_data:/data
    healthcheck:
      test: ["CMD", "redis-cli", "ping"]
      interval: 5s
```

```

timeout: 3s
retries: 10

kafka:
  image: bitnami/kafka:3.7
  container_name: ep_kafka
  ports:
    - "9092:9092"
  environment:
    - KAFKA_CFG_NODE_ID=1
    - KAFKA_CFG_PROCESS_ROLES=broker,controller
    - KAFKA_CFG_CONTROLLER_QUORUM_VOTERS=1@kafka:9093
    - KAFKA_CFG_LISTENERS=PLAINTEXT://:9092,CONTROLLER://:9093
    - KAFKA_CFG_ADVERTISED_LISTENERS=PLAINTEXT://localhost:9092
    -
    - KAFKA_CFG_LISTENER_SECURITY_PROTOCOL_MAP=PLAINTEXT:PLAINTEXT,CONTROLLER:PLAINTEXT
    - KAFKA_CFG_CONTROLLER_LISTENER_NAMES=CONTROLLER
    - KAFKA_CFG_INTER_BROKER_LISTENER_NAME=PLAINTEXT
    - ALLOW_PLAINTEXT_LISTENER=yes
    - KAFKA_CFG_AUTO_CREATE_TOPICS_ENABLE=false
  healthcheck:
    test: ["CMD-SHELL", "kafka-topics.sh --bootstrap-server localhost:9092 --list 1>/dev/null"]
    interval: 10s
    timeout: 10s
    retries: 20

kafka-init:
  image: bitnami/kafka:3.7
  container_name: ep_kafka_init
  depends_on:
    kafka:
      condition: service_healthy
  volumes:
    - ./kafka-init/create-topics.sh:/create-topics.sh:ro
  entrypoint: ["/bin/bash", "/create-topics.sh"]

volumes:
  ep_postgres_data:
  ep_redis_data:

```

Create: [infra/kafka-init/create-topics.sh](#)

```

#!/bin/bash
set -e

echo "Creating Kafka topics..."

kafka-topics.sh --bootstrap-server kafka:9092 --create --if-not-exists \
--topic events.raw --partitions 6 --replication-factor 1

kafka-topics.sh --bootstrap-server kafka:9092 --create --if-not-exists \
--topic events.dlq --partitions 3 --replication-factor 1

echo "Topics created."

```

**What this does:**

- Runs Postgres with a persistent volume (`ep_postgres_data`)
  - Runs Redis with append-only persistence (optional but nice)
  - Runs Kafka in KRaft mode (no Zookeeper)
  - Creates topics explicitly (auto-create disabled)
- 

## 3) Data Contracts (Event Schema)

### Ingestion API input payload (producer sends)

**POST** `/events`

Headers:

- `X-API-Key: <tenant key>`
- `Idempotency-Key: <unique key per logical event>` (required for dedupe)
- `Content-Type: application/json`

Body (minimal):

```
{
  "event_id": "evt_123",
  "event_type": "user_login",
  "schema_version": 1,
  "occurred_at": "2026-01-02T20:00:00Z",
```

```
"payload": { "user_id": "u1", "ip": "1.2.3.4" }  
}
```

## What gets published to Kafka (`events.raw`)

Wrap raw event with ingestion metadata (in Go before publishing):

```
{  
  "tenant_id": "tenant_a",  
  "received_at": "2026-01-02T20:00:01Z",  
  "request_id": "req_uuid",  
  "idempotency_key": "idem_abc",  
  "event": { ...original event... }  
}
```

---

## 4) Storage (Postgres)

### Table `events`

Stored after validation/enrichment in processor.

Columns:

- `id` (uuid pk)
- `tenant_id` (text)
- `event_id` (text)
- `idempotency_key` (text)
- `event_type` (text)
- `schema_version` (int)
- `occurred_at` (timestamptz)
- `received_at` (timestamptz)
- `processed_at` (timestamptz)
- `payload` (jsonb)
- `status` (text) // e.g. "processed"
- indexes:
  - `(tenant_id, event_id)`

- (tenant\_id, idempotency\_key)
- (event\_type, occurred\_at)

## Optional table `dlq_events`

If you want to persist DLQ samples in DB (optional for MVP). Otherwise, Admin API can consume DLQ topic to fetch sample messages.

---

# 5) Redis Responsibilities (Ingestion Service)

Redis keys (namespaced):

## Rate limiting

Key: `r1:{tenant_id}:{minute_bucket}`

- Use INCR + EXPIRE
- Example limit: 300 req/min per tenant
- If over limit, respond `429 Too Many Requests`

## Idempotency

Key: `idem:{tenant_id}:{idempotency_key}`

- On first request: SETNX + EXPIRE (TTL 30 minutes)
- If already exists: treat as duplicate
  - Return 200/202 with `{ "duplicate": true }` OR return 409 (your choice; recommend returning 202 with duplicate flag for friendliness)

---

# 6) Services: Responsibilities & File-by-File Expectations

## A) Go Ingestion Service (`ingest-go/`)

### Config (env vars)

- `INGEST_PORT=8080`
- `REDIS_ADDR=localhost:6379`
- `KAFKA_BROKERS=localhost:9092`
- `RATE_LIMIT_PER_MIN=300`
- `IDEMPOTENCY_TTL_SECONDS=1800`
- `API_KEYS=tenant_a:key_a,tenant_b:key_b` (simple config for MVP)

### Required endpoints

1. `POST /events`
  - Auth with API key
  - Rate limit
  - Validate JSON minimally
  - Idempotency check
  - Publish to Kafka `events.raw`
  - Return 202 with `request_id` and status
2. `GET /health`
  - returns ok + dependencies basic check (optional)
3. `GET /metrics` (optional)
  - simple JSON metrics (requests, errors, latency p95 approximate) OR expose Prometheus later

### File expectations

- `cmd/ingest/main.go`: wire config, init Redis + Kafka producer, start HTTP server
- `internal/config/config.go`: load env vars, parse API keys mapping
- `internal/http/router.go`: register routes + middleware chain
- `internal/http/middleware_auth.go`: read `X-API-Key`, map to `tenant_id`, attach tenant to request context
- `internal/http/middleware_ratelimit.go`: call redis ratelimit logic, block if exceeded
- `internal/http/handler_events.go`: parse request, call idempotency, publish to kafka, respond
- `internal/redis/client.go`: redis connection init
- `internal/redis/ratelimit.go`: INCR/EXPIRE logic
- `internal/redis/idempotency.go`: SETNX + TTL

- `internal/kafka/producer.go`: producer init + `PublishRawEvent(...)`
- `internal/http/models.go`: request/response structs
- `internal/http/responses.go`: consistent JSON error format
- `internal/metrics/metrics.go`: measure latency per request, counts (in-memory is fine)

## Response formats

Success (new event):

```
{ "status": "accepted", "request_id": "req_uuid", "duplicate": false }
```

Duplicate:

```
{ "status": "accepted", "request_id": "req_uuid", "duplicate": true }
```

Errors (standard):

```
{ "error": { "code": "RATE_LIMITED", "message": "Too many requests" } }
```

---

## B) Java Processor Service (`processor-java/`)

### Config (`application.yml`)

- Kafka bootstrap: `localhost:9092`
- topic names: `events.raw`, `events.dlq`
- Postgres: host `localhost`, db `event_platform`, user/pass from compose
- consumer group: `event-processor`
- retry policy: basic (MVP: 0-1 retry then DLQ)

### Core behavior

Consume `events.raw`:

1. Deserialize wrapper
2. Validate required fields + schema\_version allowed
3. Enrich: set processed\_at timestamp
4. Insert into Postgres

- If fails validation/deserialization → publish original message + reason to `events.dlq`

## Admin endpoints (can live in same Spring Boot app)

- `GET /admin/overview`
  - events per minute (last 1m, 5m)
  - error counts (if tracked)
- `GET /admin/top-event-types?sinceMinutes=1440`
- `GET /admin/event/search?tenant=...&eventId=...`
- `GET /admin/event/by-idempotency?tenant=...&idempotencyKey=...`
- `GET /admin/dlq/sample?limit=20` (either read from DLQ topic or from dlq table)
- `GET /admin/health`
- `GET /admin/kafka/lag` (simple lag estimate)

## File expectations

- `kafka/RawEventConsumer.java`: `@KafkaListener(topics="events.raw")` receives message
- `service/ValidationService.java`: checks required fields & version
- `service/EnrichmentService.java`: adds processed\_at and possibly normalized fields
- `service/EventProcessorService.java`: orchestrates validate → enrich → persist; catches exceptions, routes to DLQ
- `kafka/DlqProducer.java`: publishes to `events.dlq` with reason
- `db/EventEntity.java`: JPA entity mapping for events table
- `db/EventRepository.java`: Spring Data repository queries for search endpoints
- `api/AdminController.java`: REST endpoints for dashboard
- `observability/LagService.java`: compute consumer lag (MVP: expose placeholder or use Kafka AdminClient to fetch offsets)
- `resources/db/migration/V1__init.sql`: create schema

## DLQ message format (publish reason)

```
{
  "failed_at": "2026-01-02T20:05:00Z",
  "reason": "VALIDATION_FAILED: missing event_type",
  "original": { ...raw kafka message... }
}
```

---

## C) React Dashboard (`dashboard/`)

### Pages (MVP)

1. Overview
  - cards: events last minute, events last 5 minutes, top event type today
  - health status (green/red)
2. Pipeline
  - kafka lag
  - processed events per minute (last 5m)
  - DLQ count
3. DLQ
  - list sample DLQ messages
  - show reason + expand payload
4. Search
  - search by event\_id or idempotency\_key
  - show event details from DB

### API client

- Base URL: `http://localhost:8081` (if processor/admin runs on 8081)
- `api/client.ts`: fetch wrapper
- `api/types.ts`: TS types

### File expectations

- `pages/Overview.tsx`: fetch `/admin/overview`
  - `pages/Pipeline.tsx`: fetch `/admin/kafka/lag`, `/admin/overview`
  - `pages/Dlq.tsx`: fetch `/admin/dlq/sample`
  - `pages/Search.tsx`: fetch search endpoints, render event
- 

## D) Python Load Generator (`loadgen/`)

### Purpose

Simulate producers:

- sustained traffic
- duplicates

- malformed events
- multiple API keys/tenants
- measure latency and error rate

## Behavior

- Accept CLI args: `--rps`, `--minutes`, `--duplicate-rate`, `--bad-rate`, `--tenants`
- Generate `Idempotency-Key` deterministically for duplicates
- Print summary stats: success, duplicates, rate-limited, invalid, avg latency, p95 approx

## File expectations

- `loadgen.py`: main script
  - `requirements.txt`: `requests`
- 

## 7) Ports & Local Run Plan (Windows)

- Docker services:
    - Postgres: `localhost:5432`
    - Redis: `localhost:6379`
    - Kafka: `localhost:9092`
  - Apps:
    - Go ingest: `localhost:8080`
    - Java processor/admin: `localhost:8081`
    - React dashboard: `3000` (Next)
    - Loadgen hits Go ingest on `8080`
- 

## 8) Step-by-Step Build Order (What the new LLM should do)

### Step 1 — Bring up infra

- Implement `infra/docker-compose.yml` and `create-topics.sh`

- Run: `docker compose up -d` from `infra/`
- Verify:
  - Postgres up
  - Redis ping
  - Kafka topics exist

## Step 2 — Go ingest minimal

- `POST /events` publishes wrapper to Kafka `events.raw`
- No rate limit/idempotency yet

## Step 3 — Java processor consumes and writes Postgres

- Create migrations and events table
- Consume raw events and insert
- Send malformed to DLQ

## Step 4 — Add Redis rate limiting

- per-tenant limit with INCR/EXPIRE

## Step 5 — Add Redis idempotency

- SETNX with TTL
- duplicates return accepted+duplicate flag

## Step 6 — Admin endpoints

- Add overview queries from Postgres
- Add search endpoints
- Add DLQ sample (basic)

## Step 7 — React dashboard

- Build pages to show overview, pipeline, dlq, search

## Step 8 — Loadgen + failure demo

- Python script load test

- Simulate stopping processor and observing lag/DLQ
- 

## 9) “Definition of Done” (MVP Acceptance Criteria)

1. `POST /events` requires API key and idempotency key
  2. Rate limiting returns 429 for over-quota tenants
  3. Duplicate idempotency keys do not republish to Kafka
  4. Processor stores valid events in Postgres
  5. Invalid events end up in Kafka `events.dlq` (and show in DLQ page)
  6. Admin API returns:
    - events/min last 1 and 5 minutes
    - top event types today
    - search by event\_id/idempotency
    - DLQ sample list
  7. Dashboard renders those endpoints
  8. Loadgen can demonstrate:
    - duplicates ignored
    - malformed goes DLQ
    - sustained throughput
    - rate limiting works
- 

## 10) Notes / Constraints (avoid scope creep)

- No “exactly once” claims
  - Kafka provides durability + ordering per partition, not global ordering
  - Idempotency is “best effort” with Redis TTL (good enough for MVP)
  - Keep schemas simple; just enforce required fields + version
- 

Also generate: a **README.md** that explains how to run everything