
Project: Internal Developer Event Platform (Ingestion + Processing + Admin Ops)

Goal

Build an internal, developer-facing event platform with:

- **Fast ingestion API** (Go) with **API-key auth, rate limiting, idempotency**
 - **Async pipeline via Kafka** (`events.raw`, `events.dlq`)
 - **Processor** (Java Spring Boot) that validates/enriches events and stores them in **Postgres**
 - **DLQ routing for poison messages**
 - **Admin API** (can live in processor service) exposing operational metrics
 - **React admin dashboard** consuming Admin API
 - **Python load generator** simulating producers (normal, duplicates, malformed, multi-tenant)
 - Everything runs locally via **Docker Compose** for infra (Kafka/Redis/Postgres)
-

1) Repo Structure (Monorepo)

event-platform/

 README.md

 .gitignore

infra/

 docker-compose.yml

 kafka-init/

 create-topics.sh

ingest-go/

 go.mod

 cmd/

 ingest/

 main.go

 internal/

```
config/
  config.go
http/
  router.go
  middleware_auth.go
  middleware_ratelimit.go
  handler_events.go
  models.go
  responses.go
redis/
  client.go
  ratelimit.go
  idempotency.go
kafka/
  producer.go
metrics/
  metrics.go
util/
  time.go
  uuid.go
Dockerfile (optional later)
```

```
processor-java/
  pom.xml
  src/main/java/com/yourorg/processor/
    ProcessorApplication.java
  config/
    KafkaConfig.java
    PostgresConfig.java
    MetricsConfig.java
  kafka/
    RawEventConsumer.java
    DlqProducer.java
    EventPublisher.java (optional)
  service/
    EventProcessorService.java
    ValidationService.java
    EnrichmentService.java
  db/
    EventEntity.java
    EventRepository.java
    Migrations.md (notes)
  api/
    AdminController.java
```

```
DtoModels.java  
observability/  
    LagService.java  
    HealthIndicators.java  
util/  
    JsonUtil.java  
src/main/resources/  
    application.yml  
db/migration/  
    V1_init.sql  
Dockerfile (optional later)  
  
dashboard/  
    package.json  
    vite.config.ts (or next.config.js if Next.js)  
src/  
    main.tsx  
    App.tsx  
    api/  
        client.ts  
        types.ts  
    pages/  
        Overview.tsx  
        Pipeline.tsx  
        Dlq.tsx  
        Search.tsx  
components/  
    MetricCard.tsx  
    Chart.tsx (optional)  
    Table.tsx  
styles/  
    index.css  
  
loadgen/  
    requirements.txt  
loadgen.py  
scenarios/  
    baseline.json  
    spike.json
```

Key point: For MVP, you can run services directly on Windows terminals (Go/Java/React/Python) while Docker runs Kafka/Redis/Postgres.

2) Infrastructure: Docker Compose (Kafka + Redis + Postgres)

Create: `infra/docker-compose.yml`

```
services:
  postgres:
    image: postgres:16
    container_name: ep_postgres
    environment:
      POSTGRES_USER: ep_user
      POSTGRES_PASSWORD: ep_pass
      POSTGRES_DB: event_platform
    ports:
      - "5432:5432"
    volumes:
      - ep_postgres_data:/var/lib/postgresql/data
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U ep_user -d event_platform"]
      interval: 5s
      timeout: 5s
      retries: 10

  redis:
    image: redis:7
    container_name: ep_redis
    ports:
      - "6379:6379"
    command: ["redis-server", "--appendonly", "yes"]
    volumes:
      - ep_redis_data:/data
    healthcheck:
      test: ["CMD", "redis-cli", "ping"]
      interval: 5s
      timeout: 3s
      retries: 10

  kafka:
    image: bitnami/kafka:3.7
```

```

container_name: ep_kafka
ports:
- "9092:9092"
environment:
- KAFKA_CFG_NODE_ID=1
- KAFKA_CFG_PROCESS_ROLES=broker,controller
- KAFKA_CFG_CONTROLLER_QUORUM_VOTERS=1@kafka:9093
- KAFKA_CFG_LISTENERS=PLAINTEXT://:9092,CONTROLLER://:9093
- KAFKA_CFG_ADVERTISED_LISTENERS=PLAINTEXT://localhost:9092
-
KAFKA_CFG_LISTENER_SECURITY_PROTOCOL_MAP=PLAINTEXT:PLAINTEXT,CONTRO
LLER:PLAINTEXT
- KAFKA_CFG_CONTROLLER_LISTENER_NAMES=CONTROLLER
- KAFKA_CFG_INTER_BROKER_LISTENER_NAME=PLAINTEXT
- ALLOW_PLAINTEXT_LISTENER=yes
- KAFKA_CFG_AUTO_CREATE_TOPICS_ENABLE=false
healthcheck:
test: ["CMD-SHELL", "kafka-topics.sh --bootstrap-server localhost:9092 --list 1>/dev/null"]
interval: 10s
timeout: 10s
retries: 20

kafka-init:
image: bitnami/kafka:3.7
container_name: ep_kafka_init
depends_on:
kafka:
  condition: service_healthy
volumes:
- ./kafka-init/create-topics.sh:/create-topics.sh:ro
entrypoint: ["/bin/bash", "/create-topics.sh"]

volumes:
ep_postgres_data:
ep_redis_data:

```

Create: [infra/kafka-init/create-topics.sh](#)

```

#!/bin/bash
set -e

echo "Creating Kafka topics..."

```

```
kafka-topics.sh --bootstrap-server kafka:9092 --create --if-not-exists \
--topic events.raw --partitions 6 --replication-factor 1

kafka-topics.sh --bootstrap-server kafka:9092 --create --if-not-exists \
--topic events.dlq --partitions 3 --replication-factor 1

echo "Topics created."
```

What this does:

- Runs Postgres with a persistent volume (`ep_postgres_data`)
 - Runs Redis with append-only persistence (optional but nice)
 - Runs Kafka in KRaft mode (no Zookeeper)
 - Creates topics explicitly (auto-create disabled)
-

3) Data Contracts (Event Schema)

Ingestion API input payload (producer sends)

POST `/events`

Headers:

- `X-API-Key: <tenant key>`
- `Idempotency-Key: <unique key per logical event>` (required for dedupe)
- `Content-Type: application/json`

Body (minimal):

```
{
  "event_id": "evt_123",
  "event_type": "user_login",
  "schema_version": 1,
  "occurred_at": "2026-01-02T20:00:00Z",
  "payload": { "user_id": "u1", "ip": "1.2.3.4" }
}
```

What gets published to Kafka (`events.raw`)

Wrap raw event with ingestion metadata (in Go before publishing):

```
{  
    "tenant_id": "tenant_a",  
    "received_at": "2026-01-02T20:00:01Z",  
    "request_id": "req_uuid",  
    "idempotency_key": "idem_abc",  
    "event": { ...original event... }  
}
```

4) Storage (Postgres)

Table `events`

Stored after validation/enrichment in processor.

Columns:

- `id` (uuid pk)
- `tenant_id` (text)
- `event_id` (text)
- `idempotency_key` (text)
- `event_type` (text)
- `schema_version` (int)
- `occurred_at` (timestamptz)
- `received_at` (timestamptz)
- `processed_at` (timestamptz)
- `payload` (jsonb)
- `status` (text) // e.g. "processed"
- indexes:
 - `(tenant_id, event_id)`
 - `(tenant_id, idempotency_key)`
 - `(event_type, occurred_at)`

Optional table `dlq_events`

If you want to persist DLQ samples in DB (optional for MVP). Otherwise, Admin API can consume DLQ topic to fetch sample messages.

5) Redis Responsibilities (Ingestion Service)

Redis keys (namespaced):

Rate limiting

Key: `r1:{tenant_id}:{minute_bucket}`

- Use INCR + EXPIRE
- Example limit: 300 req/min per tenant
- If over limit, respond 429 Too Many Requests

Idempotency

Key: `idem:{tenant_id}:{idempotency_key}`

- On first request: SETNX + EXPIRE (TTL 30 minutes)
- If already exists: treat as duplicate
 - Return 200/202 with { "duplicate": true } OR return 409 (your choice; recommend returning 202 with duplicate flag for friendliness)

6) Services: Responsibilities & File-by-File Expectations

A) Go Ingest Service (`ingest-go/`)

Config (env vars)

- INGEST_PORT=8080
- REDIS_ADDR=localhost:6379
- KAFKA_BROKERS=localhost:9092
- RATE_LIMIT_PER_MIN=300
- IDEMPOTENCY_TTL_SECONDS=1800
- API_KEYS=tenant_a:key_a,tenant_b:key_b (simple config for MVP)

Required endpoints

1. POST /events
 - Auth with API key
 - Rate limit
 - Validate JSON minimally
 - Idempotency check
 - Publish to Kafka events.raw
 - Return 202 with request_id and status
2. GET /health
 - returns ok + dependencies basic check (optional)
3. GET /metrics (optional)
 - simple JSON metrics (requests, errors, latency p95 approximate) OR expose Prometheus later

File expectations

- cmd/ingest/main.go: wire config, init Redis + Kafka producer, start HTTP server
- internal/config/config.go: load env vars, parse API keys mapping
- internal/http/router.go: register routes + middleware chain
- internal/http/middleware_auth.go: read X-API-Key, map to tenant_id, attach tenant to request context
- internal/http/middleware_ratelimit.go: call redis ratelimit logic, block if exceeded
- internal/http/handler_events.go: parse request, call idempotency, publish to kafka, respond
- internal/redis/client.go: redis connection init
- internal/redis/ratelimit.go: INCR/EXPIRE logic
- internal/redis/idempotency.go: SETNX + TTL
- internal/kafka/producer.go: producer init + PublishRawEvent(...)
- internal/http/models.go: request/response structs

- `internal/http/responses.go`: consistent JSON error format
- `internal/metrics/metrics.go`: measure latency per request, counts (in-memory is fine)

Response formats

Success (new event):

```
{ "status": "accepted", "request_id": "req_uuid", "duplicate": false }
```

Duplicate:

```
{ "status": "accepted", "request_id": "req_uuid", "duplicate": true }
```

Errors (standard):

```
{ "error": { "code": "RATE_LIMITED", "message": "Too many requests" } }
```

B) Java Processor Service (`processor-java/`)

Config (`application.yml`)

- Kafka bootstrap: `localhost:9092`
- topic names: `events.raw`, `events.dlq`
- Postgres: host `localhost`, db `event_platform`, user/pass from compose
- consumer group: `event-processor`
- retry policy: basic (MVP: 0-1 retry then DLQ)

Core behavior

Consume `events.raw`:

1. Deserialize wrapper
2. Validate required fields + schema_version allowed
3. Enrich: set processed_at timestamp
4. Insert into Postgres
5. If fails validation/deserialization → publish original message + reason to `events.dlq`

Admin endpoints (can live in same Spring Boot app)

- `GET /admin/overview`
 - events per minute (last 1m, 5m)
 - error counts (if tracked)
- `GET /admin/top-event-types?sinceMinutes=1440`
- `GET /admin/event/search?tenant=...&eventId=...`
- `GET /admin/event/by-idempotency?tenant=...&idempotencyKey=...`
- `GET /admin/dlq/sample?limit=20` (either read from DLQ topic or from dlq table)
- `GET /admin/health`
- `GET /admin/kafka/lag` (simple lag estimate)

File expectations

- `kafka/RawEventConsumer.java`: `@KafkaListener(topics="events.raw")` receives message
- `service/ValidationService.java`: checks required fields & version
- `service/EnrichmentService.java`: adds processed_at and possibly normalized fields
- `service/EventProcessorService.java`: orchestrates validate → enrich → persist; catches exceptions, routes to DLQ
- `kafka/DlqProducer.java`: publishes to `events.dlq` with reason
- `db/EventEntity.java`: JPA entity mapping for events table
- `db/EventRepository.java`: Spring Data repository queries for search endpoints
- `api/AdminController.java`: REST endpoints for dashboard
- `observability/LagService.java`: compute consumer lag (MVP: expose placeholder or use Kafka AdminClient to fetch offsets)
- `resources/db/migration/V1__init.sql`: create schema

DLQ message format (publish reason)

```
{  
  "failed_at": "2026-01-02T20:05:00Z",  
  "reason": "VALIDATION_FAILED: missing event_type",  
  "original": { ...raw kafka message... }  
}
```

C) React Dashboard (`dashboard/`)

Pages (MVP)

1. Overview
 - cards: events last minute, events last 5 minutes, top event type today
 - health status (green/red)
2. Pipeline
 - kafka lag
 - processed events per minute (last 5m)
 - DLQ count
3. DLQ
 - list sample DLQ messages
 - show reason + expand payload
4. Search
 - search by event_id or idempotency_key
 - show event details from DB

API client

- Base URL: `http://localhost:8081` (if processor/admin runs on 8081)
- `api/client.ts`: fetch wrapper
- `api/types.ts`: TS types

File expectations

- `pages/Overview.tsx`: fetch `/admin/overview`
 - `pages/Pipeline.tsx`: fetch `/admin/kafka/lag`, `/admin/overview`
 - `pages/Dlq.tsx`: fetch `/admin/dlq/sample`
 - `pages/Search.tsx`: fetch search endpoints, render event
-

D) Python Load Generator (`loadgen/`)

Purpose

Simulate producers:

- sustained traffic
- duplicates
- malformed events
- multiple API keys/tenants
- measure latency and error rate

Behavior

- Accept CLI args: `--rps`, `--minutes`, `--duplicate-rate`, `--bad-rate`, `--tenants`
- Generate **Idempotency-Key** deterministically for duplicates
- Print summary stats: success, duplicates, rate-limited, invalid, avg latency, p95 approx

File expectations

- `loadgen.py`: main script
 - `requirements.txt`: `requests`
-

7) Ports & Local Run Plan (Windows)

- Docker services:
 - Postgres: `localhost:5432`
 - Redis: `localhost:6379`
 - Kafka: `localhost:9092`
- Apps:
 - Go ingest: `localhost:8080`
 - Java processor/admin: `localhost:8081`
 - React dashboard: `localhost:5173` (Vite) or `3000` (Next)
 - Loadgen hits Go ingest on `8080`

8) Step-by-Step Build Order (What the new LLM should do)

Step 1 — Bring up infra

- Implement `infra/docker-compose.yml` and `create-topics.sh`
- Run: `docker compose up -d` from `infra/`
- Verify:
 - Postgres up
 - Redis ping

- Kafka topics exist

Step 2 — Go ingest minimal

- POST /events publishes wrapper to Kafka events.raw
- No rate limit/idempotency yet

Step 3 — Java processor consumes and writes Postgres

- Create migrations and events table
- Consume raw events and insert
- Send malformed to DLQ

Step 4 — Add Redis rate limiting

- per-tenant limit with INCR/EXPIRE

Step 5 — Add Redis idempotency

- SETNX with TTL
- duplicates return accepted+duplicate flag

Step 6 — Admin endpoints

- Add overview queries from Postgres
- Add search endpoints
- Add DLQ sample (basic)

Step 7 — React dashboard

- Build pages to show overview, pipeline, dlq, search

Step 8 — Loadgen + failure demo

- Python script load test
 - Simulate stopping processor and observing lag/DLQ
-

9) “Definition of Done” (MVP Acceptance Criteria)

1. `POST /events` requires API key and idempotency key
 2. Rate limiting returns 429 for over-quota tenants
 3. Duplicate idempotency keys do not republish to Kafka
 4. Processor stores valid events in Postgres
 5. Invalid events end up in Kafka `events.dlq` (and show in DLQ page)
 6. Admin API returns:
 - events/min last 1 and 5 minutes
 - top event types today
 - search by event_id/idempotency
 - DLQ sample list
 7. Dashboard renders those endpoints
 8. Loadgen can demonstrate:
 - duplicates ignored
 - malformed goes DLQ
 - sustained throughput
 - rate limiting works
-

10) Notes / Constraints (avoid scope creep)

- No “exactly once” claims
 - Kafka provides durability + ordering per partition, not global ordering
 - Idempotency is “best effort” with Redis TTL (good enough for MVP)
 - Keep schemas simple; just enforce required fields + version
-